

Programming Languages: Imperative Program Construction

11. Separation Logic I

Shin-Cheng Mu

Autumn Term, 2021

Separation Matters

- Our reasoning so far is based on an important assumption: variables, having different names, are independent from each other.
- With **var** a, b , for example, mutating a does not change the value of b .
- Otherwise most of our reasoning would fail.

Remark: Procedure Calls

- Problem with procedures with call-by-reference variables.

```
proc swap (ref x, y : Int) =  
  x := x - y; y := x + y; x := y - x
```

- $\text{swap}(a, b)$ should swap the values of a and b — we have proved so before, haven't we?
- However, $\text{swap}(a, a)$ sets a to 0.
- Extra care is needed to handle function/procedure calls, which we unfortunately won't cover in this course.

Most materials of this lecture are adapted from Reynold's course in CMU [Rey11].

Other suggested reading materials: [Rey02] [O'H19], [O'H12].

1 Dynamic Memory Management

- Another source of possible violation is the heap memory model.
- Recall: variables declared are supposed to be located in *stacks*. (Also called a *store*).

- In the heap model, programmers can allocate blocks of memories in *heaps*.
- We can store addresses of heap cells in variables, lookup the content of a heap given the address, or deallocate a cell.

Pointer Manipulation

- A pointer is a variable that stores a memory address.
- In our setting we let $\text{Addr} = \text{Int}$, and let **nil** be a unique address.
- $p := \text{cons}(1, 2)$ — allocate two consecutive heap cells, set their values to 1 and 2, and store the address of the first cell in p .
 - One has no control what the address will be, other than that it won't be **nil**.
- $x := *e$ — look up the value stored in the cell with address e , and copy the value to variable x .
- $*e := f$ — let the value stored in cell with address e be updated to f .
- **free** e — free the cell having address
- In the last three cases the address e must have been allocated.

Example

| program | store and heap |
|----------------------------|---|
| | $s : x = 3 \wedge y = 4; h : \mathbf{emp}$ |
| $x := \mathbf{cons}(1, 2)$ | $s : x = 34 \wedge y = 4$ $h : 34 \mapsto 1, 35 \mapsto 2$ |
| $y := *x$ | $s : x = 34 \wedge y = 1$ $h : 34 \mapsto 1, 35 \mapsto 2$ |
| $*(x + 1) := 3$ | $s : x = 34 \wedge y = 1$ $h : 34 \mapsto 1, 35 \mapsto 3$ |
| $\mathbf{free}(x + 1)$ | $s : x = 34 \wedge y = 1$ $h : 34 \mapsto 1$ |

Notes:

- Apart from that **cons** does not return **nil**, the program cannot predict what address (e.g. 34) **cons** would return.
- Reading from, writing to, or deallocating an address that is not yet allocated aborts the program.
- We do not have an operator that gives you the address of variables in store (like **&** in C).

Linked Lists

- We abbreviate $i \mapsto 1$ and $i + 1 \mapsto 2$ to $i \mapsto 1, 2$.
- Assume that we represent lists in heap by linked lists.
- E.g $[1, 2, 3]$ is represented in the following heap, starting from address 34:

$34 \mapsto 1, 92$
 $60 \mapsto 3, \mathbf{nil}$
 $92 \mapsto 2, 60$.

- (We will present a more formal definition later.)

In-Place List Reversal

- If the address i represents a list XS , after executing the following program, i points to **nil** and j represents the *reverse* of XS .

```

{  $i$  represents  $XS$  }
 $j := \mathbf{nil}$ 
do  $i \neq \mathbf{nil} \rightarrow k := *(i + 1)$ 
     $*(i + 1) := j$ 
     $j, i := i, k$ 
od
{  $j$  represents reverse  $XS$  }
```

- That is, the program reverts a linked list without using additional space.

- Can we prove that it is correct?
- Not that easy..! The loop only works if i and j do not share any nodes. The loop invariant would be something like:

i represents $xs \wedge$
 j represents $ys \wedge \dots$
 i and j share only **nil**.

- Furthermore, we want to ensure that other data structure in the heap should remain unchanged. Assume that we have another linked-list k , we will need in the invariant:

i represents $xs \wedge$
 j represents $ys \wedge \dots$
 i and j share only **nil** \wedge
 k and $(i \text{ union } j)$ share only **nil**.

- We need to mention every pointer in the invariant. This does not scale well.

2 Separation Logic Basics

- *Separation logic*: a logic for describing and reasoning about heaps, in which sections of heaps are *separated* by default.
- Developed by people including Reynolds and O'Hearn in early 2000's.
- Widely adopted by industry in around 2010's [O'H19].
- Recall: assertions in Hoare logic are predicates on state space (values of variables in the store).
- Assertion in separation logic are predicates on the store *and* the heap.
- We will start with an informal description and give a more formal definition later.

Store and Heap

- A *store* is a (partial) function from variable names to values: $\text{Store} = \text{Var} \rightarrow \text{Val}$, where $\text{Val} = \text{Int} \cup \text{Bool} \cup \dots$
- A *heap* is a (partial) function from addresses to integers: $\text{Heap} = \text{Int} \rightarrow \text{Int}$ — an address is also a *Int*.

- The domain of a function f is denoted $\text{dom } f$.
- We denote $\text{dom } h_0 \cap \text{dom } h_1 = \emptyset$ by $h_0 \perp h_1$.
- Given functions h_0 and h_1 where $h_0 \perp h_1$, define

$$\begin{aligned} (h_0 \cdot h_1) x &= h_0 x \text{ if } x \in \text{dom } h_0, \\ &= h_1 x \text{ if } x \in \text{dom } h_1. \end{aligned}$$

Some Primitives

Given a heap h ,

- **emp** h holds if $\text{dom } h = \emptyset$.
 - **emp** says that nothing is allocated in the heap.
- $e \mapsto e'$ holds of h if $\text{dom } h = \{e\}$ and $h e = e'$.
 - h is a singleton heap containing only e' in address e .
 - Note that both e and e' are expressions!
- $P * Q$ holds of h if $h = h_0 \cdot h_1$ and $P h_0$ and $Q h_1$.
 - That $h_0 \cdot h_1$ being defined implies that $h_0 \perp h_1$.
 - h can be decomposed into two *disjoint* heap h_0 and h_1 such that p holds of h_0 and q holds of h_1 .
- *True* holds of any h , while *False* holds of no h .
- $e \mapsto e_0, e_1, \dots, e_n \equiv (e \mapsto e_0) * (e+1 \mapsto e_1) * \dots * (e+n \mapsto e_n)$.
- $e \mapsto _ \equiv \langle \exists v :: e \mapsto v \rangle$.
- $e \mapsto e' \equiv (e \mapsto e') * \text{True}$.
- *separating implication* $p \multimap q$ will be introduced later.

The True Story

- The presentation above was very simplified.
- In fact, all the predicates introduced above are predicate on *store and heap*, because we need the store to evaluate an expression.
- We will keep it simple for now. For a more precise account, see Reynolds [Rey02, Rey11].
- Keep in mind, for example, that $x \mapsto 3$, where x is a variable, actually means x is mapped to some a in the store, and a is mapped to 3 in the heap.
 - The predicate can be invalidated if either the value of x or the value stored in the heap changes.

Examples

- $x \mapsto 3, y$.
- $(x \mapsto 3, y) * (y \mapsto 3, x)$.
- $(x \mapsto 3, y) \wedge (y \mapsto 3, x)$.
- $(x \mapsto 3, y) \wedge (y \mapsto 3, x)$.

Separating Implication

- *Separating implication* is defined by:

$$(P \multimap Q) h = \langle \forall h_0 : h_0 \perp h \wedge P h_0 : Q (h_0 \cdot h) \rangle .$$

- That is, $P \multimap Q$ holds of h if, given any h_0 that is disjoint from h and satisfies P , we have $h_0 \cdot h$ satisfies Q .

Example

- Suppose P asserts various things, including $x \mapsto 3, 4$. Thus P holds of

$$\begin{aligned} s : x &= a \\ h : a \mapsto 3, a+1 \mapsto 4, \text{rest of heap} \end{aligned}$$

- $(x \mapsto 3, 4) \multimap P$ holds of the following store and heap:

$$\begin{aligned} s : x &= a \\ h : \text{rest of heap} \end{aligned}$$

- $(x \mapsto 1, 2) * ((x \mapsto 3, 4) \multimap P)$ holds of the following store and heap:

$$\begin{aligned} s : x &= a \\ h : a \mapsto 1, a+1 \mapsto 2, \text{rest of heap} \end{aligned}$$

Heap Mutation – Motivation

- From the example above we notice that

$$\begin{aligned} \{ &(x \mapsto 1) * ((x \mapsto 3) \multimap P) \} \\ *x &:= 3 \\ \{ &P \} \end{aligned}$$

- To be slightly more general,

$$\begin{aligned} \{ &(x \mapsto _) * ((x \mapsto 3) \multimap P) \} \\ *x &:= 3 \\ \{ &P \} \end{aligned}$$

- We will see a more general rule later.

3 Commands and Rules

Rule of Constancy

- In logic systems, the following notation denotes “ Q can be established by establishing P ”:

$$\frac{P}{Q}$$

- In Hoare logic, the following “rule of constancy” holds:

$$\frac{\{P\} S \{Q\}}{\{P \wedge R\} S \{Q \wedge R\}}$$

- It allows us to reason about programs in a more modular way.
- However, rule of constancy does not hold for programs allowing dynamic memory management. The following does not hold, for example.

$$\frac{\{x \mapsto _ \} * x := 4 \{x \mapsto 4\}}{\{x \mapsto _ \wedge y \mapsto 3\} * x := 4 \{x \mapsto 4 \wedge y \mapsto 3\}}$$

- (What if x and y evaluate to the same address?)

Frame Rule

- With the introduction of separating conjunction, we do have:

$$\frac{\{P\} S \{Q\}}{\{P * R\} S \{Q * R\}}$$

- The rule above is called the “frame rule”. With it we can again reason about programs modularly.
- Wanting to have such rule is the very reason why separation logic was developed.

Commands

- Now we discuss rules associated with each pointer manipulation command.
- Each command is associated with three types of rule: local, global (forward), and backward rules.

Mutation

- Local: $\{e \mapsto _ \} * e := e' \{e \mapsto e'\}$.
- Global: $\{(e \mapsto _) * R\} * e := e' \{(e \mapsto e') * R\}$.
- Backwards: $\{(e \mapsto _) * ((e \mapsto e') \multimap P)\} * e := e' \{P\}$.
- The global rule is often the result of applying frame rule to the local rule.

Deallocation

- Local: $\{e \mapsto _ \} \text{free } e \{\text{emp}\}$.
- Global: $\{(e \mapsto _) * R\} \text{free } e \{R\}$.
- For this case, the global rule is also a backwards rule.

Allocation, Non-Overwriting

A simpler, *non-overwriting* case, where x does not occur free in e .

- Local: $\{\text{emp}\} x := \text{cons } e \{x \mapsto e\}$,
- Global: $\{R\} x := \text{cons } e \{(x \mapsto e) * R\}$.
- Backwards rule and the general case is much more complex — we will discuss them later.
- We have not yet discussed the rule for looking up ($x := *e$) — which turns out to be surprisingly complex. Discussion postponed.

Example

The following code fragment tries to glue together adjacent cells, if possible.

```

{(x ↦ -) * (y ↦ -)}
if y = x + 1 → skip
| x = y + 1 → x := y
| |x - y| > 1 → free x; free y
               x := cons (1, 2)
fi
{x ↦ -, -}

```

Allocation, General Case

- Local:

$$\{x = X \wedge \text{emp}\} x := \text{cons } e \{x \mapsto e[x \backslash X]\} ,$$

where X is distinct from x and does not occur free in e .

- Global:

$$\{R\} x := \text{cons } e \{(\exists x_0 :: x \mapsto e[x \backslash x_0]) * R[x \backslash x_0]\} ,$$

where x_0 is distinct from x and does not occur free in e and R .

- Backwards:

$$\{(\forall x_1 :: (x_1 \mapsto e) \multimap P[x \backslash x_1])\} x := \text{cons } e \{P\} ,$$

where x_1 is distinct from x and does not occur free in e and R .

Lookup, Non-Overwriting

Provided that x does not occur free in e ,

- Local: $\{e \mapsto v\} x := *e \{x = v \wedge e \mapsto x\}$.
- Global:

$$\{\langle \exists v :: (e \mapsto v) * R[x \setminus v] \rangle\} x := *e \{(e \mapsto x) * R\} ,$$

where v does not occur free in e and R .

Lookup, General

- Local:

$$\{x = x_0 \wedge e \mapsto v\} x := *e \{x = v \wedge e[x \setminus x_0] \mapsto x\} ,$$

where x, x_0, v distinct.

- Global:

$$\begin{aligned} & \{\langle \exists v :: (e \mapsto v) * R[x_0 \setminus x] \rangle\} \\ & x := *e \\ & \{\langle \exists x_0 :: (e[x \setminus x_0] \mapsto x) * R[v \setminus x] \rangle\} , \end{aligned}$$

where x, x_0, v distinct, x_0 and v not free in e , x not free in R .

- Backwards:

$$\begin{aligned} & \{\langle \exists v :: (e \mapsto v) * ((e \mapsto v) \multimap P[x \setminus v]) \rangle\} \\ & x := *e \\ & \{P\} \end{aligned}$$

- Backwards, in a shorter form:

$$\begin{aligned} & \{\langle \exists v :: (e \hookrightarrow v) \wedge P[x \setminus v] \rangle\} \\ & x := *e \\ & \{P\} \end{aligned}$$

References

- [O'H12] P. W. O'Hearn. A primer on separation Logic (and automatic program verification and analysis). In T. Nipkow, O. Grumberg, and B. Hauptmann, editors, *Software Safety and Security: Tools for Analysis and Verification*, volume 33 of *NATO Science for Peace and Security Series*, pages 286–318, 2012.
- [O'H19] P. W. O'Hearn. Separation logic. *Communications of the ACM*, 62(2):86–95, February 2019.
- [Rey02] J. C. Reynolds. Separation logic: a logic for shared mutable data structures. In G. D. Plotkin, editor, *Annual IEEE Symposium on Logic in Computer Science*, pages 55–74. IEEE Computer Society Press, 2002.
- [Rey11] J. C. Reynolds. 15-818A3 Introduction to Separation Logic. Carnegie Mellon University. <https://www.cs.cmu.edu/afs/cs.cmu.edu/project/fox-19/member/jcr/www15818As2011/cs818A3-11.html>, 2011.